

# Setup

```
import pandas as pd
import numpy as np
import os

from sklearn.utils import shuffle
from sklearn.model_selection import train_test_split,
GroupShuffleSplit
from sklearn.metrics import confusion_matrix, accuracy_score,
f1_score, precision_score, recall_score

import tensorflow as tf
from tensorflow import math
from tensorflow.keras import layers, models, metrics, optimizers,
Input, Model
from tensorflow.keras.utils import plot_model
from tensorflow.keras import backend as K

from pandas import DataFrame
import seaborn as sn

try:
    from livelossplot import PlotLossesKeras
except ImportError:
    ! pip install -q -U livelossplot
    from livelossplot import PlotLossesKeras

from tqdm import tqdm
from matplotlib import pyplot as plt
from matplotlib.pyplot import figure

try:
    import keras_tuner as kt
except ImportError:
    ! pip install -q -U keras-tuner
    import keras_tuner as kt

===== 0.0/128.9 kB ? eta -:--:--
===== 41.0/128.9 kB 876.2 kB/s eta
0:00:01 ===== 128.9/128.9 kB 1.8
MB/s eta 0:00:00

def plot_conf_matrix(true_labels, predicted_labels):
    cm = confusion_matrix(true_labels, predicted_labels) # rows=true,
    cols=pred
    sum_true = tf.expand_dims(tf.reduce_sum(cm, axis=1), axis=1) # sum
    per row (total # true)
```

```

norm_cm = (cm / sum_true).numpy()

df = DataFrame(norm_cm, index=ALL_LABELS, columns=ALL_LABELS)
# df[df == 0] = np.nan # to hide cell if value=0

fig, ax = plt.subplots(figsize=(5,5))
heatmap = sn.heatmap(df, cmap='Blues', cbar=False, square=True,
annot=True, fmt='.1%', ax=ax, linewidths=.2, linecolor=[0,.1,.2])
ax.tick_params(left=False, bottom=False)
plt.xlabel("Predicted y")
plt.ylabel("True y")
plt.show()

```

## Dataset

```

"""
Each class is represented by a number
0      : airplane
1      : automobile
2      : bird
3      : cat
4      : deer
5      : dog
6      : frog
7      : horse
8      : ship
9      : truck
"""
NUM_LABELS = 10

(x_train, y_train), (x_test, y_test) =
tf.keras.datasets.cifar10.load_data()
y_train = y_train.reshape((50000,))
y_test = y_test.reshape((10000,))
assert x_train.shape == (50000, 32, 32, 3)
assert x_test.shape == (10000, 32, 32, 3)
assert y_train.shape == (50000,)
assert y_test.shape == (10000,)

x_train_norm = x_train / 255.0
x_test_norm = x_test / 255.0

Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-
python.tar.gz
170498071/170498071 [=====] - 31s 0us/step

```

```
ALL_LABELS = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog',  
'frog', 'horse', 'ship', 'truck']  
labels_dict = {i:ALL_LABELS[i] for i in range(10)}
```

# Model

## Creation

```
hp = {  
    'dense1_len': 65,  
    'dense2_len': 77,  
    'learning_rate': 1e-4,  
  
    'dropout_rate': 0.33 ,  
  
    'conv1_nfilters': 67,  
    'conv1_kernel_size': 3,  
    'conv1_strides': 1,  
  
    'maxpool1_width': 2,  
    'maxpool1_stride': 2,  
  
    'conv2_nfilters': 12,  
    'conv2_kernel_size': 4,  
    'conv2_strides': 1,  
  
    'maxpool2_width': 2,  
    'maxpool2_stride': 2,  
  
    'conv3_nfilters': 23,  
    'conv3_kernel_size': 4,  
    'conv3_strides': 1,  
  
    'maxpool3_width': 2,  
    'maxpool3_stride': 2,  
}  
  
def get_model(x_shape, hp):  
    x = Input(name='in_x', shape=x_shape)  
  
    ...  
  
    Conv blocks  
    ...  
  
    conv = layers.Conv2D(filters=hp['conv1_nfilters'],  
kernel_size=hp['conv1_kernel_size'], strides=hp['conv1_strides'],  
                        name='conv1', activation='relu',  
kernel_regularizer='l2', padding='same')(x)
```

```

conv = layers.BatchNormalization()(conv)
conv = layers.Dropout(name='dropout1', rate=hp['dropout_rate'])
(conv)
conv = layers.MaxPooling2D(name='maxpool1',
pool_size=hp['maxpool1_width'], strides=hp['maxpool1_stride'])(conv)

```

```

conv = layers.Conv2D(filters=hp['conv2_nfilters'],
kernel_size=hp['conv2_kernel_size'], strides=hp['conv2_strides'],
name='conv2', activation='relu',
kernel_regularizer='l2', padding='same')(x)
conv = layers.BatchNormalization()(conv)
conv = layers.Dropout(name='dropout2', rate=hp['dropout_rate'])
(conv)
conv = layers.MaxPooling2D(name='maxpool2',
pool_size=hp['maxpool2_width'], strides=hp['maxpool2_stride'])(conv)

```

'''  
*After adding the second block, make sure to complete the fully connected part below. It's similar to what we did yesterday. I've already flattened it for you.*  
'''

*Now, train your model. See if it does well. Let me know once you've done this.*

*After training with two blocks and talking to me, try adding a third "block" to improve performance.*  
'''

```

'''
Fully connected
'''
flatten = layers.Flatten(name='flatten')(conv)
fc = layers.Dense(name='fc1', units=hp['dense1_len'],
activation='relu', kernel_regularizer='l2')(flatten)
fc = layers.Dense(name='fc2', units=hp['dense2_len'],
activation='relu', kernel_regularizer='l2')(fc)

```

```

out = layers.Dense(name='out', activation='softmax',
units=NUM_LABELS)(fc)

```

```

# compiling the model...
model = Model(name="model", inputs=x, outputs=out)
model.compile(

optimizer=optimizers.Adam(learning_rate=hp['learning_rate'],
clipnorm=1.0), # clipnorm=1.0 does gradient clipping to stop gradient

```

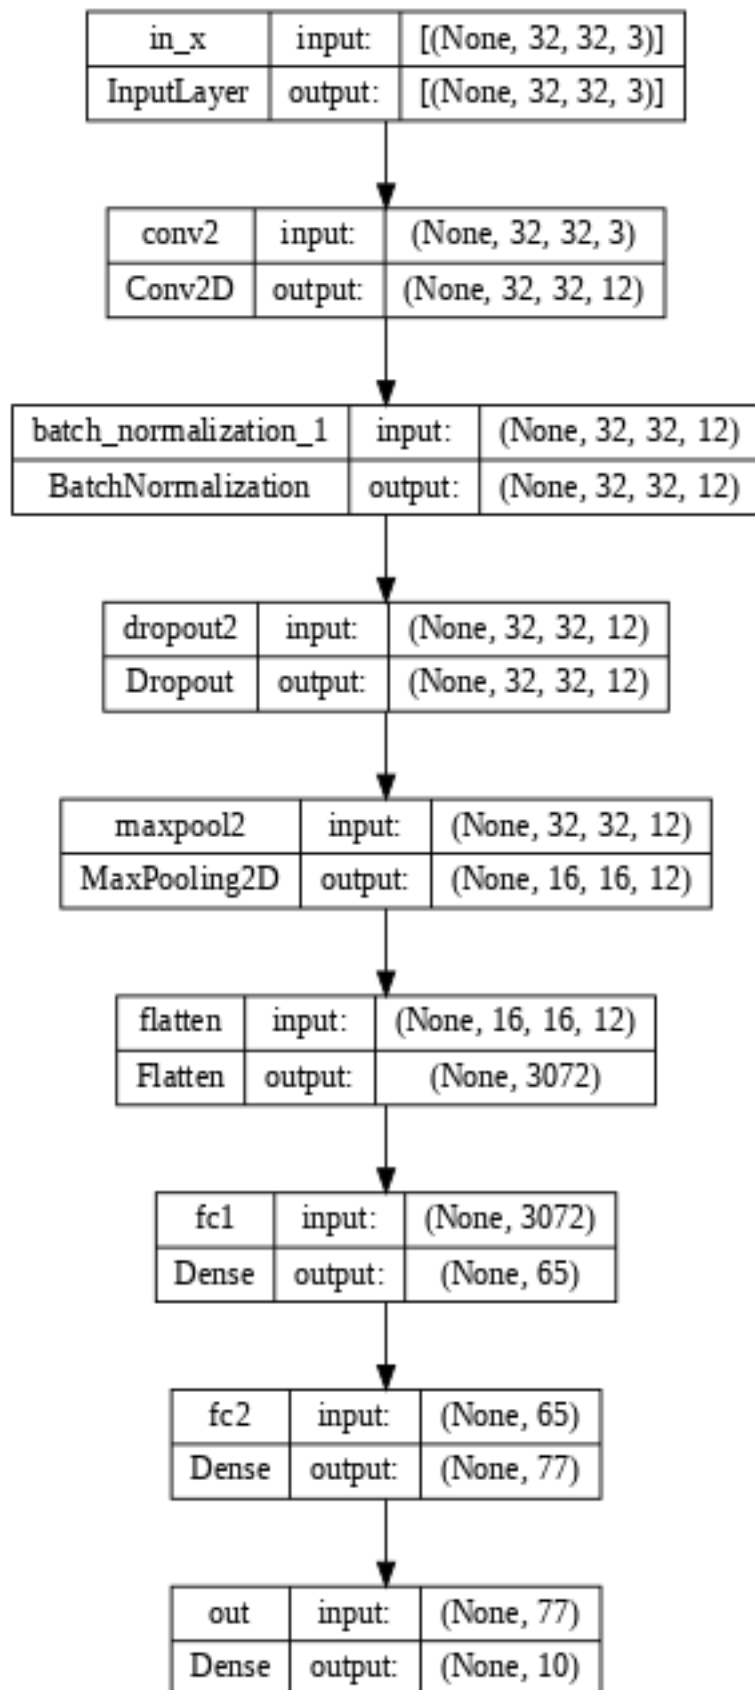
*explosion*

```
        loss=tf.keras.losses.CategoricalCrossentropy(),  
        metrics=['accuracy'])  
    return model
```

```
model = get_model(x_train_norm[0].shape, hp)  
model.summary()  
plot_model(model, show_shapes=True, dpi=64)
```

Model: "model"


Layer (type)	Output Shape	Param #
in_x (InputLayer)	[(None, 32, 32, 3)]	0
conv2 (Conv2D)	(None, 32, 32, 12)	588
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 12)	48
dropout2 (Dropout)	(None, 32, 32, 12)	0
maxpool2 (MaxPooling2D)	(None, 16, 16, 12)	0
flatten (Flatten)	(None, 3072)	0
fc1 (Dense)	(None, 65)	199745
fc2 (Dense)	(None, 77)	5082
out (Dense)	(None, 10)	780
Total params: 206,243		
Trainable params: 206,219		
Non-trainable params: 24		



# Training

```
tf.random.set_seed(42)
model = get_model(x_train_norm[0].shape, hp)
history = model.fit(x_train_norm, tf.one_hot(y_train,
depth=NUM_LABELS),
                    epochs=40,

callbacks=[tf.keras.callbacks.EarlyStopping(monitor='val_loss',
patience=10),
          PlotLossesKeras()],
          validation_data=(x_test_norm,
tf.one_hot(y_test, depth=NUM_LABELS))
)
```



```
accuracy
  training      (min:   0.396, max:   0.783, cur:
0.783)
  validation    (min:   0.422, max:   0.654, cur:
0.646)
Loss
  training      (min:   0.905, max:   3.128, cur:
0.905)
  validation    (min:   1.248, max:   2.685, cur:
```

```
1.280)
1563/1563 [=====] - 8s 5ms/step - loss:
0.9051 - accuracy: 0.7826 - val_loss: 1.2803 - val_accuracy: 0.6461
```

## Results

### Training Performance

```
y_train_probs = model.predict(x_train_norm)
model_train_pred = [np.argmax(y) for y in y_train_probs]

-----
-----
NameError                                Traceback (most recent call
last)
<ipython-input-2-c58d7949fb16> in <cell line: 1>()
----> 1 y_train_probs = model.predict(x_train_norm)
      2 model_train_pred = [np.argmax(y) for y in y_train_probs]

NameError: name 'model' is not defined

print("Training Performance")
print("Accuracy", accuracy_score(y_train, model_train_pred))
print("F1 Score (macro)", f1_score(y_train, model_train_pred,
average='macro'))
plot_conf_matrix(y_train, model_train_pred)
```

### Test Performance

```
y_test_probs = model.predict(x_test_norm)
model_test_pred = [np.argmax(y) for y in y_test_probs]

print("Test Performance")
print("Accuracy", accuracy_score(y_test, model_test_pred))
print("F1 Score (macro)", f1_score(y_test, model_test_pred,
average='macro'))
plot_conf_matrix(y_test, model_test_pred)

for i in range(10):
    print(f'TEST IMAGE {i}')
    print('Predicted:', labels_dict[model_test_pred[i]])
    print('Actual:', labels_dict[y_test[i]])
    print('Image:')
    plt.imshow(x_test[i], cmap='gray')
    plt.show()
```